

P4 - Concurrency Control (Part 1)

Critical Region

In exercise P3, you probably used an array to create space for each thread to store its partial sum.

If array elements happen to share a cache line, this leads to false sharing.

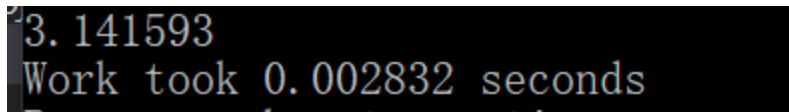
- Non-shared data in the same cache line so each update invalidates the cache line ... in essence “sloshing independent data” back and forth between threads.

The critical construct is a directive that contains a structured block. The construct allows only a single thread at a time to execute the structured block (region). Multiple critical regions may exist in a parallel region, and may act cooperatively (only one thread at a time in all critical regions), or separately (only one thread at a time in each critical region when a unique name is supplied on each critical construct). An optional (lock) hint clause may be specified on a named critical construct to provide the OpenMP runtime guidance in selection of a locking mechanism.

Q1. Modify your “pi program” from [P4Q1](#) to avoid false sharing due to the sum array by using the following steps:

1. Create a scalar local to each thread to accumulate partial sums.
2. No array, so no false sharing.
3. Sum goes “out of scope” beyond the parallel region ... so you must sum it in here. Must protect summation into pi in a critical region so updates don’t conflict.
(#pragma omp critical)

Output:



```
3.141593
Work took 0.002832 seconds
```

shared Specifies that one or more variables should be shared among all threads.

[OpenMP Clauses | Microsoft Docs](#) - #pragma omp parallel shared(var)

critical Specifies that code can only be executed on one thread at a time.

[OpenMP Directives | Microsoft Docs](#) - #pragma omp critical(var)

P4 - Concurrency Control (Part 1)

- Q2. Compute the average of array A with a vector of MAX = 100000 in parallel. Below is the serial program which you can download [here](#):

```
double ave=0.0, A[MAX];
int i;
for (i=0;i<MAX;i++){
    ave += A[i];
}
ave = ave/MAX;
```

1st Method: #pragma omp parallel, #pragma omp critical

Reference:

<https://docs.microsoft.com/en-us/cpp/parallel/openmp/reference/openmp-directives?view=msvc-170#critical>

2nd Method: #pragma omp parallel for reduction(+:ave)

Question (discussion):

Why the sequential (without #pragma omp critical) run faster / same speed as running the code in threads using #pragma omp critical ?

P4 - Concurrency Control (Part 1)

Producer Consumer

Q3. Parallelize the [P4Q3.c](#) program. This is a well known pattern called the producer consumer pattern. Using `#pragma omp sections` and `#pragma omp flush`

- One thread produces values that another thread consumes.
- Often used with a stream of produced values to implement “pipeline parallelism”.
- The key is to implement pairwise synchronization between threads.
- Producer needs to tell consumer that the data is ready
- Consumer needs to wait until data is ready
- Producer and consumer need a way to communicate data
 - output of producer is input to consumer
 - Producer and consumer often communicate through First-in-first-out (FIFO) queue

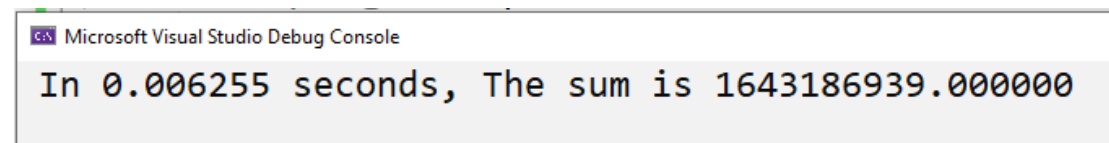
`#pragma omp sections` - Identifies code sections to be divided among all threads.

<https://docs.microsoft.com/en-us/cpp/parallel/openmp/reference/openmp-directives?view=msvc-170#sections-openmp>

`#pragma omp flush` - Specifies that all threads have the same view of memory for all shared objects.

<https://docs.microsoft.com/en-us/cpp/parallel/openmp/reference/openmp-directives?view=msvc-170#flush-openmp>

Output:



```
Microsoft Visual Studio Debug Console

In 0.006255 seconds, The sum is 1643186939.000000
```

P4 - Concurrency Control (Part 1)

The *atomic* Construct

The following example avoids race conditions (simultaneous updates of an element of x by multiple threads) by using the atomic construct.

Atomic = ~~one at a time~~ in one go

The advantage of using the atomic construct in this example is that it allows updates of two different elements of x to occur in parallel. If a critical construct were used instead, then all updates to elements of x would be executed serially (though not in any guaranteed order).

Note that the atomic directive applies only to the statement immediately following it.

As a result, elements of y are not updated atomically in this example.

- Q4. Modify the [P4Q4.c](#) and add atomic construct (`#pragma omp atomic`) inside the parallel construct to ensure the respective final sum of array of x and y are consistent.

```
The sum is 49994696.000000 and 99985792.000000
```

`#pragma omp atomic` - Specifies a memory location that will be updated atomically.

<https://docs.microsoft.com/en-us/cpp/parallel/openmp/reference/openmp-directives?view=msvc-170#atomic>