

## P5 - Concurrency Control (Part 2)

## Revision

- Q1 There are two companies with 10,000,000 employees. Both companies have the same owner. You would like to know the salaries of all employees of each company. The serial code can be downloaded [here](#). Try to execute the code in parallel by adding the directives below:

```
#pragma      omp      parallel      shared(salaries1,      salaries2)
num_threads(16)

#pragma omp for reduction(+:salaries1)

#pragma omp for reduction(+:salaries2)
```

**Challenge :** Try to make the output to be printed properly, in line by line  
(right) \* tips: use the `#pragma omp` command that we've learned before.

- Q2** Modify the solution for Q1, so that it displays the salaries once only. (Hint: use **only one line** of reduction clause)

```
Salaries1: 1362004000  
Salaries2: 1464690000
```

- Q3 Modify the solution for Q1, so that it displays the salaries once only. (Hint: use #pragma omp single)

```
Salaries1: 1362004000  
Salaries2: 1464690000
```

## O1 - OMP - Atom1C - Fur Reduction - Single.CPP

Q1 There are two companies with 10,000,000 employees. Both companies have the same owner. You would like to know the salaries of all employees of each company. The serial code can be downloaded [here](#). Try to execute the code in parallel by adding the directives below:

```
#pragma omp parallel shared(salaries1, salaries2) num_threads(16)

#pragma omp for reduction(+:salaries1)

#pragma omp for reduction(+:salaries2)
```

```
Salaries1: 1362004000  
Salaries1: 1362004000Salaries1: 1362004000  
  
Salaries2: 1464690000  
Salaries2: 1464690000  
Salaries2: 1464690000  
Salaries2: Salaries2: 1464690000  
Salaries2: 1464690000  
Salaries2: 1464690000  
Salaries2: 1464690000  
1464690000
```

**Challenge :** Try to make the output to be printed properly, in line by line  
(right) \* tips: use the `#pragma omp` command that we've learned before.

Critical

## Solution

```
int company_salaries_q1()
{
    int salaries1 = 0;
    int salaries2 = 0;
    double runtime;
    runtime = omp_get_wtime(); ↙
#pragma omp parallel shared(salaries1, salaries2) num_threads(16)
    {
#pragma omp for reduction(+ : salaries1) ←
        for (int employee = 0; employee < MAX_EMPLOYEE; employee++)
            salaries1 += fetchTheSalary(employee, Co::Company1);

#pragma omp critical ←
    {
        std::cout << "Salaries1: " << salaries1 << std::endl;
        // printf("Salaries1: %d\n", salaries1); // printf is thread safe
    }

#pragma omp for reduction(+ : salaries2) ←
    for (int employee = 0; employee < MAX_EMPLOYEE; employee++)
        salaries2 += fetchTheSalary(employee, Co::Company2);

#pragma omp critical ←
    {
        std::cout << "Salaries2: " << salaries2 << std::endl;
        // printf("Salaries2: %d\n", salaries2);
    }
    runtime = omp_get_wtime() - runtime;
    std::cout << " In " << runtime << " seconds" << std::endl;
    return 0;
}
```

## Initial Setup

## challenge:

↓ avoid race

## Condition

Q2 Modify the solution for Q1, so that it displays the salaries once only. (Hint: use **only one line of** reduction clause)

```
Salaries1: 1362004000  
Salaries2: 1464690000  
Total runtime = 1.0000000000000002
```

### Solution

```
int company_salaries_q2()  
{  
    int salaries1 = 0;  
    int salaries2 = 0;  
    double runtime;  
    runtime = omp_get_wtime();  
#pragma omp parallel shared(salaries1, salaries2) num_threads(16)  
    {  
#pragma omp for reduction(+ : salaries1, salaries2)  
        for (int employee = 0; employee < MAX_EMPLOYEE; employee++)  
        {  
            salaries1 += fetchTheSalary(employee, Co::Company1);  
            salaries2 += fetchTheSalary(employee, Co::Company2);  
        }  
        std::cout << "Salaries1: " << salaries1 << std::endl;  
        std::cout << "Salaries2: " << salaries2 << std::endl;  
    }  
    runtime = omp_get_wtime() - runtime;  
    std::cout << " In " << runtime << " seconds" << std::endl;  
    return 0;  
}
```

- Q3 Modify the solution for Q1, so that it displays the salaries once only. (Hint: use #pragma omp single)

```
:Salaries1: 1362004000  
:Salaries2: 1464690000
```

## Solution

```
int company_salaries_q3()  
{  
    int salaries1 = 0;  
    int salaries2 = 0;  
    double runtime;  
    runtime = omp_get_wtime();  
#pragma omp parallel shared(salaries1, salaries2) num_threads(16)  
    {  
#pragma omp for reduction(+ : salaries1, salaries2)  
        for (int employee = 0; employee < MAX_EMPLOYEE; employee++)  
        {  
            salaries1 += fetchTheSalary(employee, Co::Company1);  
            salaries2 += fetchTheSalary(employee, Co::Company2);  
        }  
#pragma omp single  
    {  
        std::cout << "Salaries1: " << salaries1 << std::endl;  
        std::cout << "Salaries2: " << salaries2 << std::endl;  
    }  
    }  
    runtime = omp_get_wtime() - runtime;  
    std::cout << " In " << runtime << " seconds" << std::endl;  
    return 0;  
}
```

## P5 - Concurrency Control (Part 2)

---

### Barrier

A barrier defines a point in the code where all active threads will stop until all threads have arrived at that point. With this, you can guarantee that certain calculations are finished. For instance, in this code snippet, computation of `y` can not proceed until another thread has computed its value of `x`.

```
#pragma omp parallel
{
    int mytid = omp_get_thread_num();
    x[mytid] = some_calculation();
    y[mytid] = x[mytid]+x[mytid+1];
}
```

This can be guaranteed with a barrier pragma:

```
#pragma omp parallel
{
    int mytid = omp_get_thread_num();
    x[mytid] = some_calculation();
    #pragma omp barrier
    y[mytid] = x[mytid]+x[mytid+1];
}
```

Apart from the barrier directive, which inserts an explicit barrier, OpenMP has **implicit barriers** after a load sharing construct. Thus the following code is well defined:

```
#pragma omp parallel
{
    #pragma omp for
    for (int mytid=0; mytid<number_of_threads; mytid++)
        x[mytid] = some_calculation();
    #pragma omp for
    for (int mytid=0; mytid<number_of_threads-1; mytid++)
        y[mytid] = x[mytid]+x[mytid+1];
}
```

## 02\_BARRIER.cpp

```
#include <iostream>
#include <omp.h>

using namespace std;

int some_calculation()
{
    return 10;
}

const int MAX_THREADS = 16;

void barrier_exercise()
{
    int x[MAX_THREADS] = {0}, y[MAX_THREADS] = {0};

#pragma omp parallel num_threads(MAX_THREADS)
{
    int mytid = omp_get_thread_num();
    x[mytid] = some_calculation();

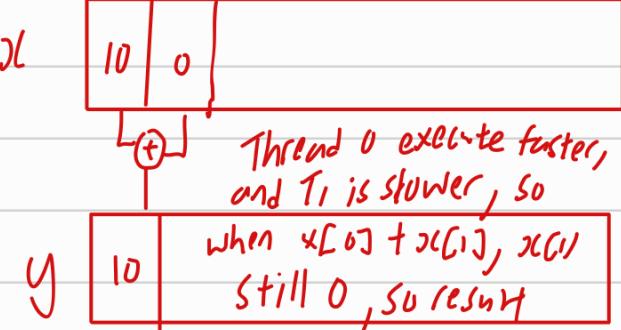
#pragma omp barrier
    y[mytid] = x[mytid] + x[mytid + 1];
    for (int i = 0; i < MAX_THREADS; i++)
        cout << "y[" << i << "] = " << y[i] << endl;
}

int main()
{
    barrier_exercise();
    return 0;
}
```

→ if don't have barrier

T<sub>0</sub>    T<sub>1</sub>

10	0
----	---



10, not 20

```
y[0] = 20
y[1] = 20
y[2] = 20
y[3] = 20
y[4] = 20
y[5] = 20
y[6] = 20
y[7] = 20
y[8] = 20
y[9] = 20
y[10] = 20
y[11] = 20
y[12] = 20
y[13] = 20
y[14] = 20
y[15] = 30
```

after use barrier

x	10   10   10   10   10   10   10   ....
---	---

barrier  
ensure x all done, then go to y

y	20   20   20   20   20   20   20   ....
---	---

```
y[0] = 10
y[1] = 10
y[2] = 10
y[3] = 20
y[4] = 10
y[5] = 20
y[6] = 20
y[7] = 10
y[8] = 20
y[9] = 10
y[10] = 20
y[11] = 20
y[12] = 10
y[13] = 10
y[14] = 20
y[15] = 10
```

### Implicit Barrier

At the end of a parallel region the team of threads is dissolved and only the master thread continues. Therefore, there is an implicit barrier at the end of a parallel region.

```
#pragma omp parallel
{
    some_calculation();
} // Implicit barrier on every end of a parallel region
```

There is some barrier behavior associated with `omp for` loops and other work sharing constructs barriers. For instance, there is an implicit barrier at the end of the loop. This barrier behavior can be canceled with the clause.

You will often see the idiom

```
#pragma omp parallel
{
    #pragma omp for nowait
    for (i=0; i<N; i++)
        a[i] = // some expression
    #pragma omp for
    for (i=0; i<N; i++)
        b[i] = ..... a[i] .....
```

*cancel out*

*implicit block  
(wait the for loop complete  
then continue execute other  
line)*

Here the `nowait` clause implies that threads can start on the second loop while other threads are still working on the first. Since the two loops use the same schedule here, an iteration that uses `a[i]` can indeed rely on that value has been computed.

### Pros and Cons

The main reason for a barrier in a program is to avoid data races and to ensure the correctness of the program.

Of course there are some downsides. Each synchronization is a threat for performance. When a thread **waits** for other threads, it does not do any useful work and it spends valuable resources.

Another problem might occur if we are not carefully inserting barriers. As soon as one thread reaches the barrier then all threads in the team must reach the barrier. Otherwise, the threads waiting at the barrier will **wait forever**.

## P5 - Concurrency Control (Part 2)

---

### Locks

OpenMP also has the traditional mechanism of a lock. A lock is somewhat similar to a critical section: it guarantees that some instructions can only be performed by one process at a time. However, a critical section is indeed about code; a lock is about **data**. With a lock you make sure that some data elements can only be touched by one process at a time.

One simple example of the use of locks is generation of a histogram . A histogram consists of a number of bins that get updated depending on some data. Here is the basic structure of such a code:

```
int count[100];
float x = some_function();
int ix = (int)x;
if (ix >= 100)
    error();
else
    count[ix]++;
```

It would be possible to guard the last line:

```
#pragma omp critical
count[ix]++;
```

but that is unnecessarily restrictive. If there are enough bins in the histogram, and if the `some_function` takes enough time, there are unlikely to be conflicting writes. The solution then is to create an array of locks, with one lock for each count location.

Create/destroy:

```
void omp_init_lock(omp_lock_t *lock);
void omp_destroy_lock(omp_lock_t *lock);
```

Set and release:

```
void omp_set_lock(omp_lock_t *lock);
void omp_unset_lock(omp_lock_t *lock);
```

Since the set call is blocking, there is also

```
omp_test_lock(omp_lock_t *lock);
```

Unsetting a lock needs to be done by the thread that set it. Lock operations implicitly have a **flush**.

#pragma omp critical

{  
CR/CS  
}

#pragma omp atomic

↑  
arithmetic  
statement

✗ omp-lock\_t lk;

OMP-init\_lock(&lk);

⋮

OMP-set-lock(&lk);

[  
C.R./C.S.

OMP-unset-lock(&lk);

OMP-destroy-lock(&lk);

## LOCK

✓ flexibly can unlock and  
lock in a if else  
case

✗ difficult to use

✗ error prone

if()

lock()

else()

unlock()

Thread

Library

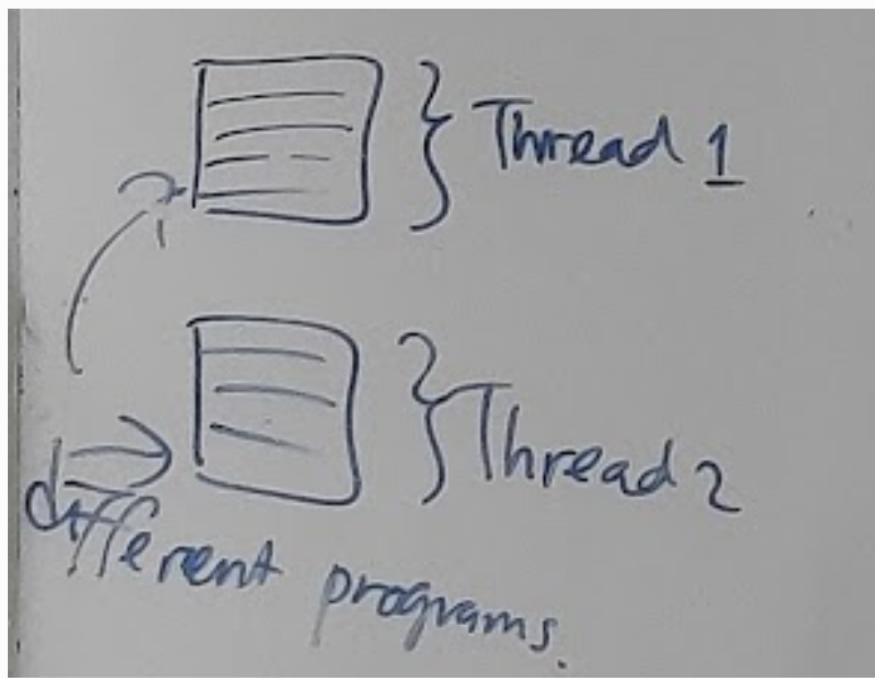
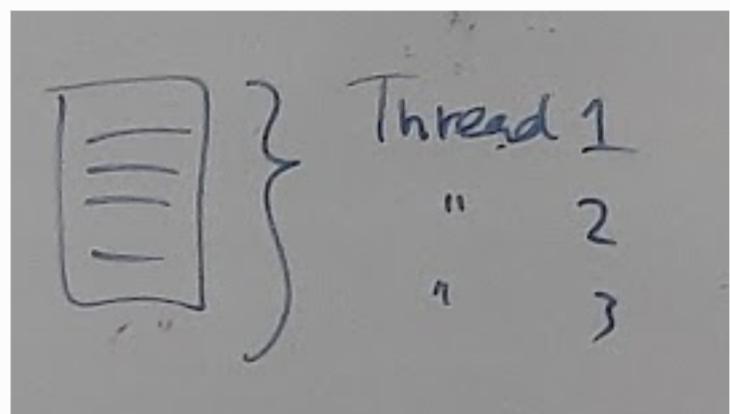
- ↳ concurrent computing
- ↳ NO guarantee to use multiple cores

C/C++

- ↳ cannot expect performance gain

OMP

- ↳ parallel computing
- ↳ expect performance gain due to OMP use multiple cores.



assign different  
thread on  
some  
program

assign diff thread on  
diff program

## P5 - Concurrency Control (Part 2)

Besides using OpenMP locks, the C++ Thread class can be utilized to perform parallel concurrent processing using multiple threads.

### Thread

Class to represent individual threads of execution.

<https://m.cplusplus.com/reference/thread/thread/>

Example of code with Threads process

```
// thread example
#include <iostream>           // std::cout
#include <thread>             // std::thread

void foo() {
    // do stuff...
}

void bar(int x) {
    // do stuff...
}

int main() {
    std::thread first (foo);    // spawn new thread that calls foo()
    std::thread second (bar,0); // spawn new thread that calls bar(0)
    std::cout << "main, foo and bar now execute concurrently...\n";

    // synchronize threads:
    first.join();              // pauses until first finishes
    second.join();              // pauses until second finishes

    std::cout << "foo and bar completed.\n";
    return 0;
}
```

How many threads?

3 threads

a) Main thread

b) Main thread spawns first thread

c) Main thread spawns second thread

if donot include this, after main thread end,  
first and second thread will get destroy  
at the same time when mutex is getting  
destroy

### Mutex

A mutex is a [lockable object](#) that is designed to signal when critical sections of code need exclusive access, preventing other threads with the same protection from executing concurrently and accessing the same memory locations.

<https://m.cplusplus.com/reference/mutex/mutex/>

# 3 - Thread library

```
// thread example
#include <iostream> // std::cout
#include <thread> // std::thread
// #include <windows.h> // Sleep() function is only available in Windows
#include <chrono> // Portable, can use in both windows and Linux environment

void foo()
{
    // do stuff...
    // Sleep only available in Windows (Because from window.h)
    // Sleep(4000);
    std::chrono::milliseconds timespan(4000); ←
    std::this_thread::sleep_for(timespan);
}

void bar(int x)
{
    // do stuff...
    // Sleep(5000);
    std::chrono::milliseconds timespan(5000); ←
    std::this_thread::sleep_for(timespan);
}

int main()
{
    std::thread first(foo); // spawn new thread that calls foo()
    std::thread second(bar, 0); // spawn new thread that calls bar(0)

    std::cout << "main, foo and bar now execute concurrently...\n";

    // synchronize threads:
    first.join(); // pauses until first finishes
    second.join(); // pauses until second finishes

    std::cout << "foo and bar completed.\n";

    return 0;
}
```

Diagram annotations:

- A red bracket groups the two `std::this_thread::sleep_for(timespan);` statements, labeled "concurrent total run : 5sec".
- A red bracket groups the two `Sleep` statements, labeled "5sec".

Output:

```
main, foo and bar now execute concurrently...
foo and bar completed. ← wait 5sec
```

# 04\_Mutex.cpp

```
#include <iostream>
#include <thread>
#include <mutex>
#include <vector>

std::mutex mtx; // mutex for critical section

void print_block(int n, char c)
{
    // critical section (exclusive access to std::cout signaled by locking mtx):
    mtx.lock(); ←
    for (int i = 0; i < n; ++i)
    {
        std::cout << c;
    }
    std::cout << '\n';
    mtx.unlock(); ←
}

int main()
{
    std::thread th1(print_block, 10000, '*');
    std::thread th2(print_block, 10000, '$');

    th1.join();
    th2.join();

    return 0;
}
```

Output:

Without Mutex :

```
*****$$$$$$$$$$$$$$*****$**$*****$*$*$*$*$*$**  
*****$**$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$  
*****$*$*$*****$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$  
*$*$*****$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$  
*****$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$  
*****$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$  
***$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$  
**$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$  
*****$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$*$
```

↳ race condition

With Mutex

```
*****  
*****  
*****  
$$$$$$$$$$$$$$  
$$$$$$$$$$$$$$  
$$$$$$$$$$$$$$
```

## P5 - Concurrency Control (Part 2)

---

### Lock mutex

Locks all the objects passed as arguments, blocking the calling thread if necessary.

<https://mcplusplus.com/reference/mutex/lock/>

- Q4 Besides using OpenMP locks, the C++ <mutex> library also allows the program to initialize locks to prevent data races. Declare std::mutex increment and then implement the methods below for the [code](#) that count the number of even numbers:

```
increment.lock();
```

```
increment.unlock();
```

\* Add

```
std::cout << "thread: " << numEven << std::endl;
```

after the increment of numEven to see the difference of implementing / not implementing lock mutex.

# 05\_Mutex\_NumEven.cpp

```
#include <iostream>
#include <thread>
#include <mutex>
#include <vector>

std::mutex increment;

void countEven(const std::vector<int> &numbers,
               int &numEven)
{
    increment.lock();
    for (const auto n : numbers)
    {
        if (n % 2 == 0)
        {
            numEven++;
        }
    }
    increment.unlock();
}

void insert_numbers_into_vector(std::vector<int> &v, int start_val, int len)
{
    for (int i = start_val; i < start_val + len; i++)
        v.insert(v.end(), i);
}

int main()
{
    // std::vector<int> numbers1 = { 1, 2, 3, 4, 5, 6, 7, 8, 9, 10 };
    // std::vector<int> numbers2 = { 11, 12, 13, 14, 15, 16, 17, 18 };
    std::vector<int> numbers1;
    std::vector<int> numbers2;
    int n = 0;

    insert_numbers_into_vector(numbers1, 0, 10000);
    insert_numbers_into_vector(numbers2, 20000, 10000);

    std::thread t1(countEven, std::ref(numbers1), std::ref(n));
    std::thread t2(countEven, std::ref(numbers2), std::ref(n));

    t1.join();
    t2.join();

    std::cout << "Total even numbers is: " << n << std::endl;
    return 0;
}
```

for( int i=0; i < numbers.size(); i++ ) {  
 int n = numbers[i];  
 if( n % 2 == 0 ) {  
 numEven++;  
 }  
}

Numbers1  
[0 1 2 3 4] ... [ ]  
10,000

numbers2  
[20,000 20,001 20,002] ... [ ]  
29,999

Output.

With mutex

```
Total even numbers is: 10000
```

without mutex

```
Total even numbers is: 9914
```

by race condition

## P5 - Concurrency Control (Part 2)

---

### Lock guard

A *lock guard* is an object that manages a *mutex object* by keeping it always locked.

[https://m.cplusplus.com/reference/mutex/lock\\_guard/](https://m.cplusplus.com/reference/mutex/lock_guard/)

### Deadlock

- Q5 Download the demo [code](#) for deadlock. Implement locks at the correct section to prevent deadlock. To avoid the deadlock:
- The simplest solution is to always lock the mutexes in the same order.
  - Use individual mutex lock and unlock for any critical sections.
  - The std::lock function can lock 2 or more mutexes at once without risk of a deadlock.

Place the locks below at the correct code section.

```
// std::lock(m_print, m_bag);
// put the lock_guard in the correct order
std::lock_guard<std::mutex> printGuard(m_print,
std::adopt_lock);
std::lock_guard<std::mutex> bagGuard(m_bag, std::adopt_lock);
```

```
Making an apple juice...
I made an excellent apple juice for you.
I threw out all pears from your bag!
Bag: x x x x x x x x
Press any key to continue . . .
```

```
std::mutex m_print;
std::mutex m_bag;

void makeAppleJuice(std::vector<std::string> &bag)
{
    std::lock_guard<std::mutex> bagGuard(m_bag); ← need to follow
    std::lock_guard<std::mutex> printGuard(m_print); fixed order
    if not
    std::cout << "Making an apple juice..." << std::endl; it
    // std::lock_guard<std::mutex> bagGuard(m_bag); will
    int numApples = 0; cause
    for (std::size_t i = 0; i != bag.size(); i++) deadlock
    {
        if (bag[i] == "a")
        {
            numApples++;
            bag[i] = "x";
        }
    } ← LR for bagGuard
    if (numApples < 5)
    {
        std::cout << "I can not make an apple juice." << std::endl;
    }
    else
    {
        std::cout << "I made an excellent apple juice for you." << std::endl;
    }
}
```

```

void throwOutPear(std::vector<std::string> &bag)
{
    // std::lock_guard<std::mutex> printGuard(m_print);

    std::lock_guard<std::mutex> bagGuard(m_bag); } need to
    for (std::size_t i = 0; i != bag.size(); i++)
    {
        if (bag[i] == "p")
        {
            bag[i] = "x";
        }
    }

    std::lock_guard<std::mutex> printGuard(m_print); } in fixed
    std::cout << "I threw out all pears from your bag!" << std::endl;
}

```

need to  
in fixed  
order  
if not  
deadlock

```

void printItem(const std::string &fruit)
{
    std::cout << fruit << " ";
}

```

```

int main()
{
    std::vector<std::string> bag = {"a", "a", "a", "p", "p", "a", "p", "a"};

    std::thread t1(makeAppleJuice, std::ref(bag));
    std::thread t2(throwOutPear, std::ref(bag));
    t1.join();
    t2.join();

    std::cout << "Bag: ";
    std::for_each(bag.begin(), bag.end(), printItem);
    std::cout << std::endl;

    return 0;
}

```

What makeAppleJuice function does?

Making an apple juice...  
I made an excellent apple juice for you.  
Bag: x x x p p x p x

$\downarrow \text{numApple} > 5$

What does throwOutPear Does

I threw out all pears from your bag!  
Bag: a a a x x a x a

Output if does not follow lock order:

Deadlock

└ ./06\_Deadlock

Making an apple juice...

↳ but sometime deadlock won't happen

Making an apple juice...

I made an excellent apple juice for you.

I threw out all pears from your bag!

Bag: x x x x x x x x

If follow order guarantee, no dead lock

Making an apple juice...

I made an excellent apple juice for you.

I threw out all pears from your bag!

Bag: x x x x x x x x

Mutex mtx1, mtx2;

Thread 1

```
task1 {  
    mtx1.lock();  
    mtx2.lock();  
    do_the_critical_task1();  
    mtx1.unlock();  
    mtx2.unlock();  
}
```

Thread 2

```
task2 {  
    mtx2.lock();  
    mtx1.lock();  
    do_the_critical_task2();  
    mtx1.unlock();  
    mtx2.unlock();  
}
```

thread 1

mtx1.lock()

timeline  
↓

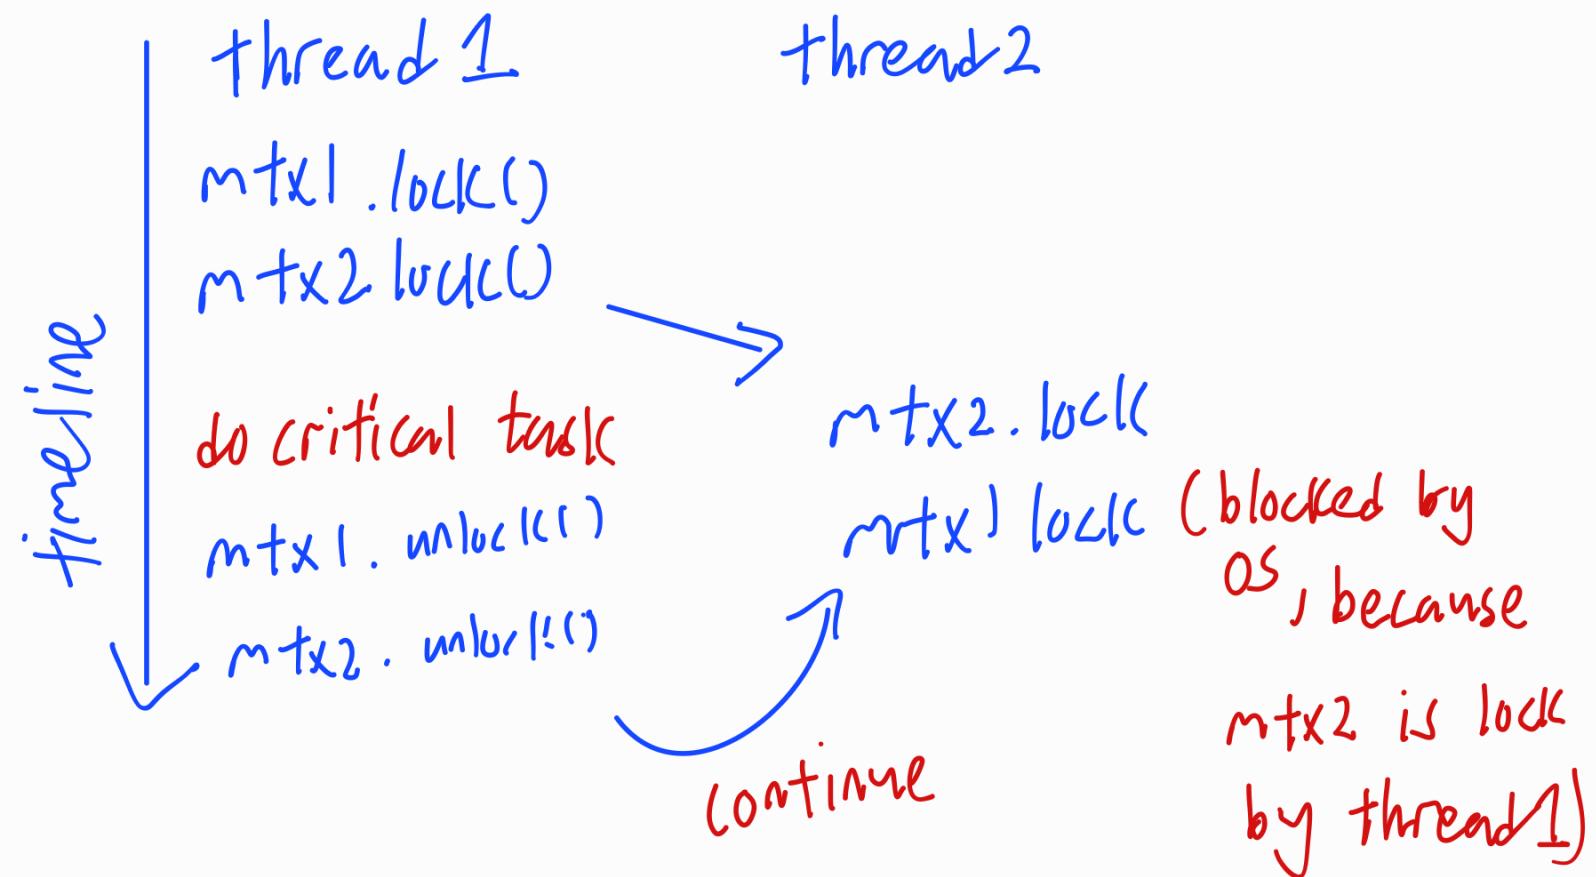
mtx2.lock()  
(blocked)

→ mtx2.lock

deadlock

thread 2

mtx1.lock (blocked by  
OS, because  
mtx1 is lock  
by thread1)



→ not always got deadlock  
 ↴ got chance happen

Mutex mtx1, mtx2j

Thread 1

```
task1 {  
    mtx1.lock();  
    mtx2.lock();  
    do_the_critical_task1();  
    mtx1.unlock();  
    mtx2.unlock();  
}
```

In order to

prevent

circular wait,

the mutex lock need to

be same

Thread 2

```
task2 {
```

```
    mtx1.lock();  
    mtx2.lock();
```

do\_the\_critical\_task2();

```
    mtx1.unlock();  
    mtx2.unlock();
```

thread 1

mtx1.lock()

timeline

mtx2.lock()



thread 2

mtx1.lock



(blocked by  
OS, because  
mtx1 is lock  
by thread1)

