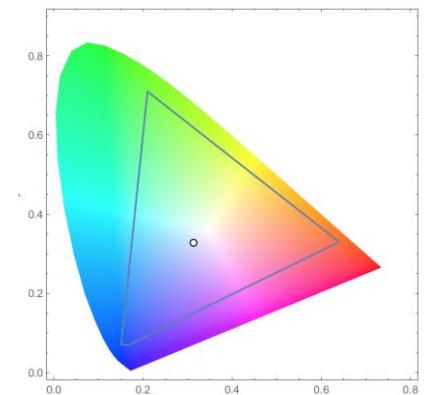
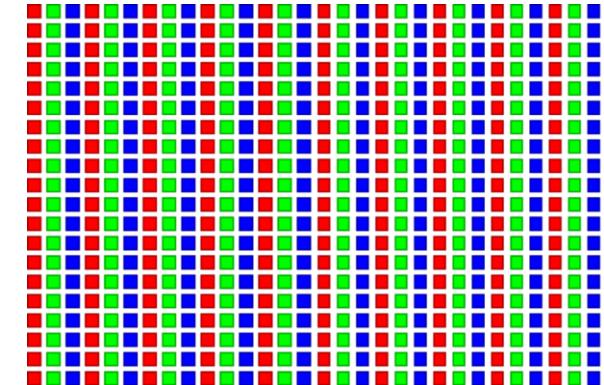


We can use RGB to represent any color

RGB COLOR IMAGE REPRESENTATION

- Each pixel in an image is an RGB value
- The format of an image's row is
 $(r\ g\ b)\ (r\ g\ b)\ \dots\ (r\ g\ b)$
- RGB ranges are not distributed uniformly
- Many different color spaces, here we show the constants to convert to AdobeRGB color space
 - The vertical axis (y value) and horizontal axis (x value) show the fraction of the pixel intensity that should be allocated to G and B. The remaining fraction ($1-y-x$) of the pixel intensity that should be assigned to R
 - The triangle contains all the representable colors in this color space



01 - To Gray Scale with Cuda

```
__global__ void to_grayscale_kernel(uint8_t *gray_img, uint8_t *bgr_img,
                                    int width, int height, int channels) {
    // Add your code here to turn the color image into grayscale image
    // Note: gray = 0.07 * blue + 0.71 * green + 0.21 * red
    // uint8_t = unsigned 8-bit integer type (Range: 0 - 255)
    // static_cast<uint8_t>(value) -> // Explicitly cast to uint8_t

    int x = threadIdx.x + blockIdx.x * blockDim.x;
    int y = threadIdx.y + blockIdx.y * blockDim.y;

    if (x < width && y < height){
        int gray_idx = y * width + x;
        int color_idx = gray_idx * channels;

        // Get BGR pixel
        uint8_t blue = bgr_img[color_idx];
        uint8_t green = bgr_img[color_idx + 1];
        uint8_t red = bgr_img[color_idx + 2];

        gray_img[gray_idx] = static_cast<uint8_t>(0.07 * blue + 0.71 * green + 0.21 * red);
    }
}
```

```
const int thread_per_blk = 32;      // Use 32 threads per block for x and y
void cuda_color_to_grayscale(uint8_t *gray_img, uint8_t *color_img,
                             int width, int height, int channels) {
    uint8_t* cData;      // Color image data for device
    uint8_t* gData;      // Grayscale image data for device

    // Compute the image data size for color and grayscale in bytes. Note that
    // the gray scale image has only one channel (one color), so the size is
    // 3 times smaller.
    const size_t grayDataSize = sizeof(uint8_t) * width * height;
    const size_t colorDataSize = grayDataSize * channels;

    // Allocate memory for gData and cData with the size of grayDataSize
    // and colorDataSize, respectively
    cudaMalloc((void**)&cData, colorDataSize);
    cudaMalloc((void**)&gData, grayDataSize);

    // Transfer BGR image from host (color_img) to device (cData). The size
    // should be colorDataSize
    cudaMemcpy(cData, color_img, colorDataSize, cudaMemcpyHostToDevice);

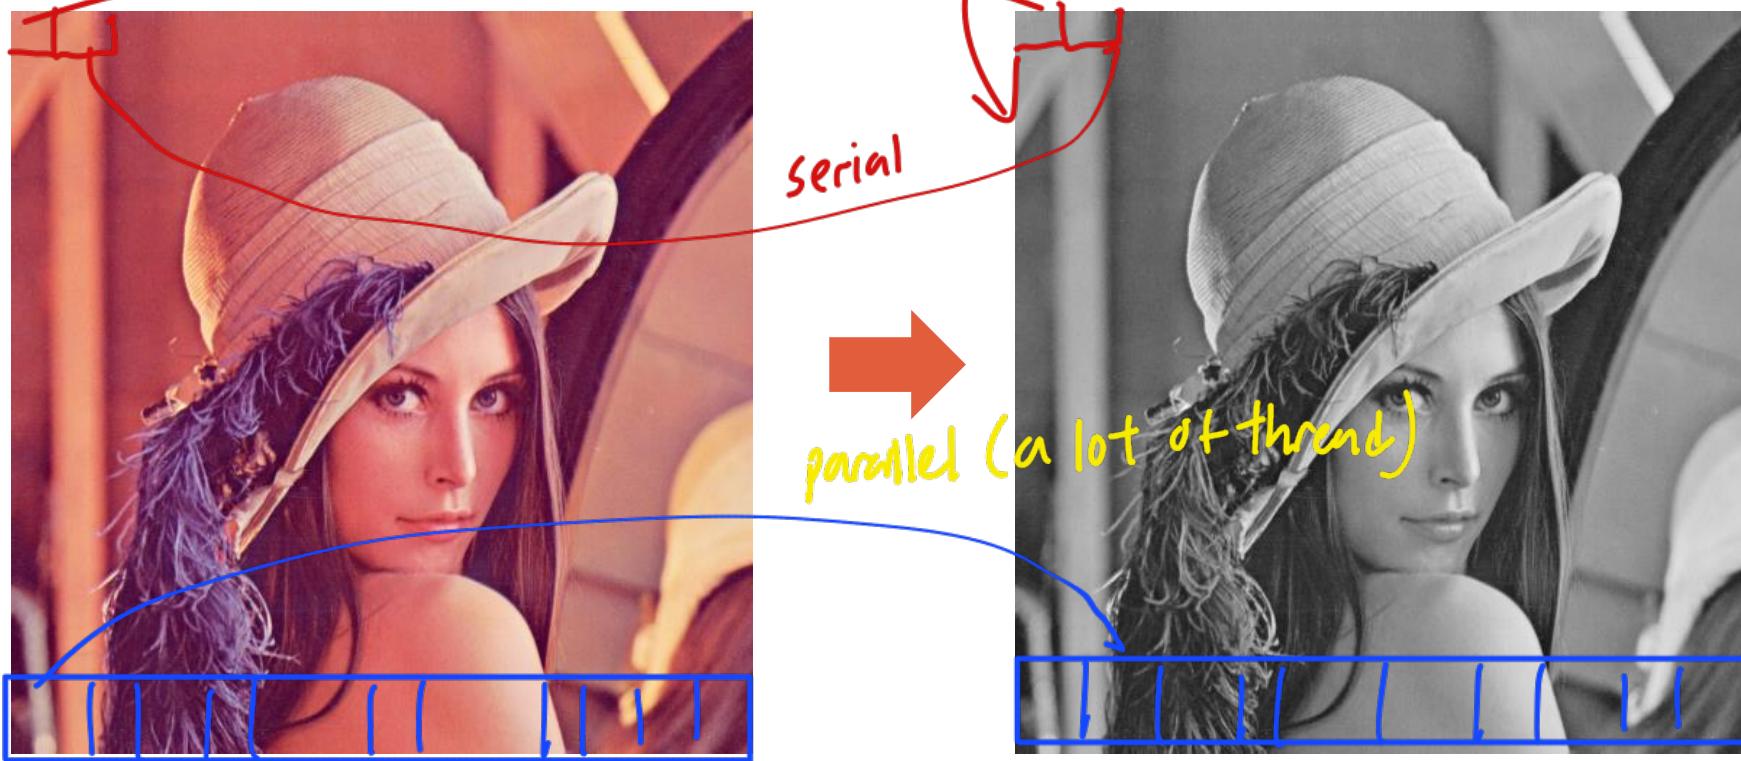
    // Calculate blocksize and gridsize, use the recommended 'thread_per_blk'
    dim3 blockSize(thread_per_blk, thread_per_blk, 1);
    dim3 gridSize((width + thread_per_blk - 1) / thread_per_blk,
                  (height + thread_per_blk - 1) / thread_per_blk);

    // Call the GPU kernel
    to_grayscale_kernel <<<gridSize, blockSize>>> (gData, cData, width, height, channels);

    // Transfer Grayscale image from device (gData) to host (gray_img). The
    // size should be colorDataSize
    cudaMemcpy(gray_img, gData, grayDataSize, cudaMemcpyDeviceToHost);

    // Free the allocated device memory
    cudaFree(cData);
    cudaFree(gData);
}
```

RGB TO GRayscale CONVERSION



A grayscale digital image is an image in which the value of each pixel carries only intensity information.

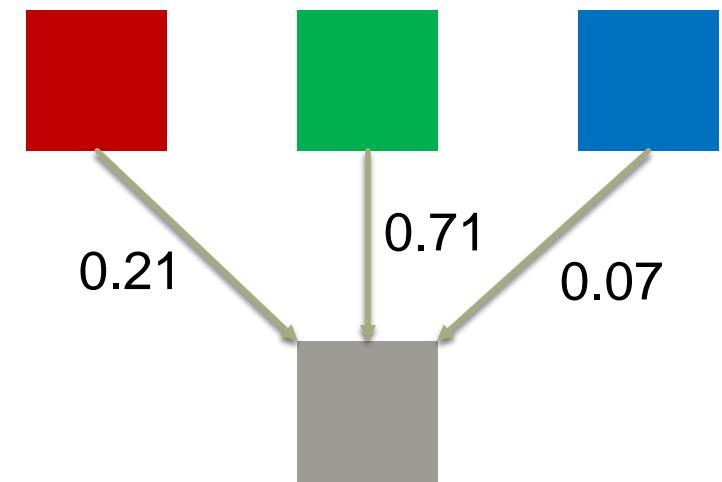
grayScale

COLOR CALCULATING FORMULA

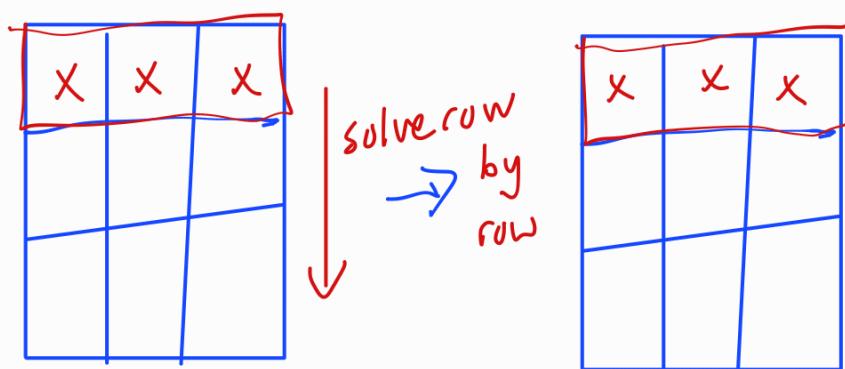
- For each pixel (r g b) at (I, J) do:

$$\text{grayPixel}[I,J] = 0.21*r + 0.71*g + 0.07*b$$

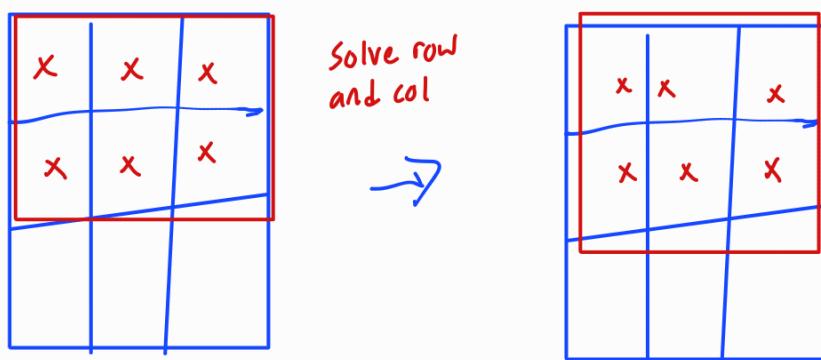
- This is just a dot product
 $\langle [r,g,b], [0.21, 0.71, 0.07] \rangle$ with the constants being specific to input RGB space



1D parallel



2D parallel



color
image

gray Scale
image

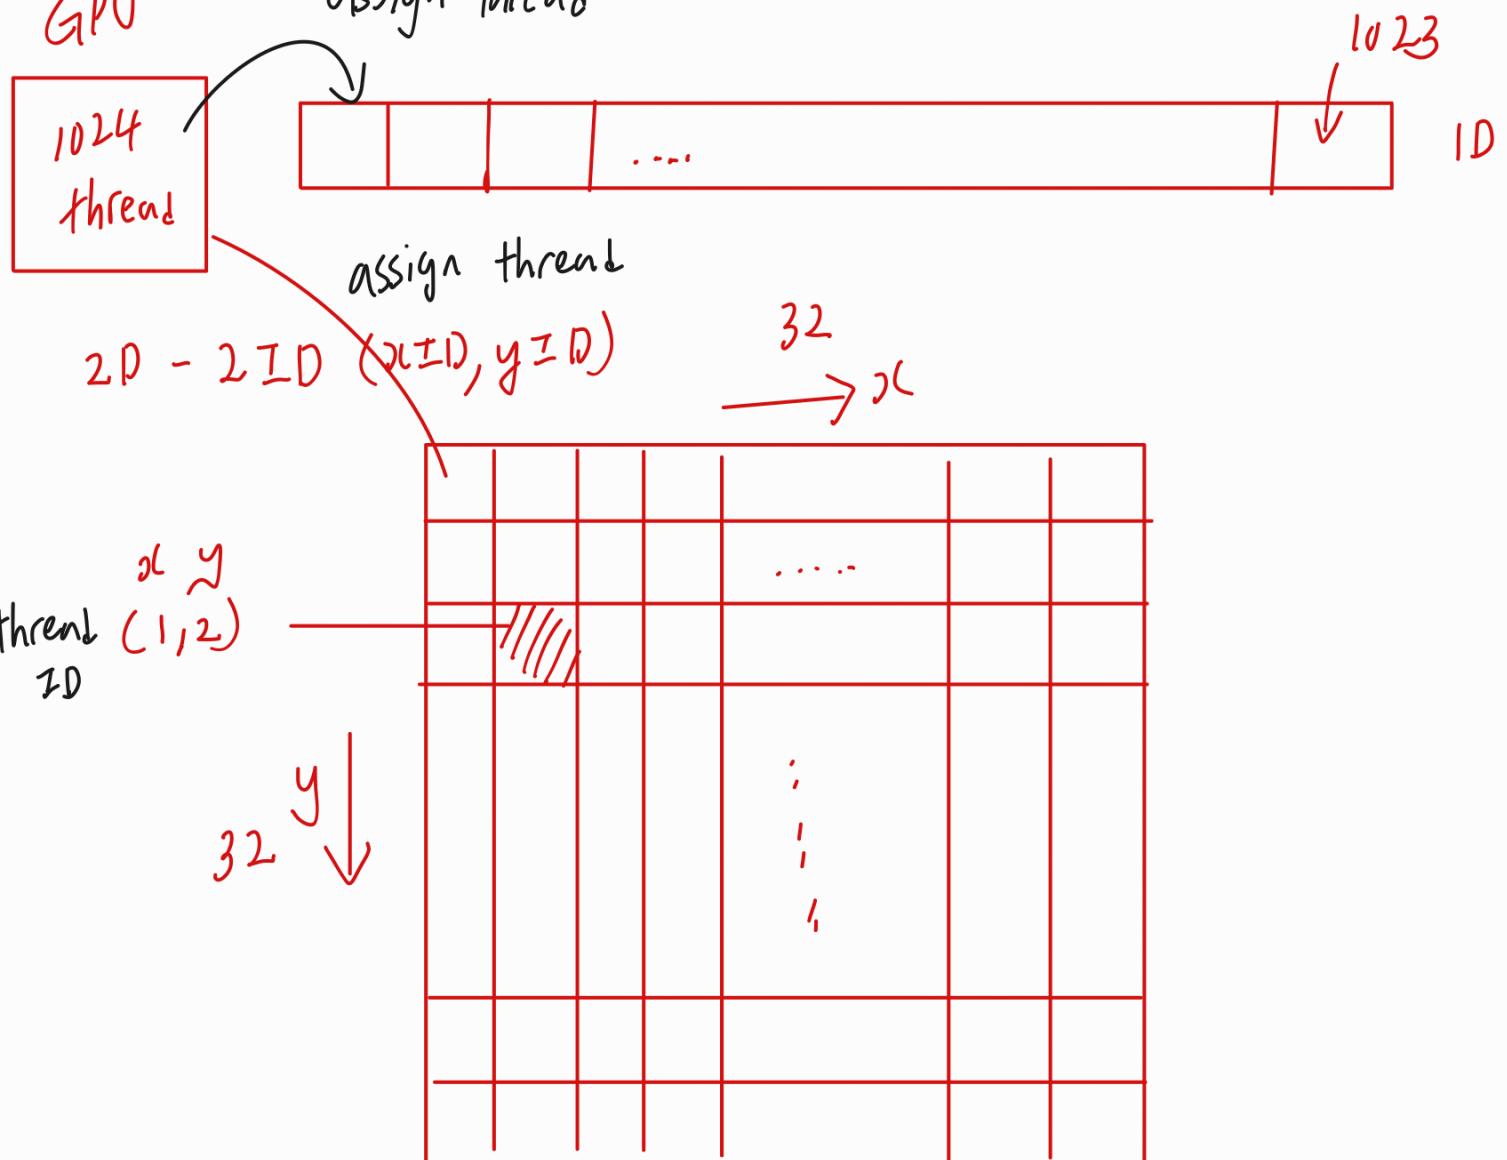
GPU

1D - 1ID (x (ID))

GPU

1024
thread

assign thread



3D

RGB TO GRayscale CONVERSION CODE

```
// we have 3 channels corresponding to RGB  
// The input image is encoded as unsigned characters [0, 255]  
__global__ void colorConvert(unsigned char * grayImage,  
                                unsigned char * rgbImage,
```

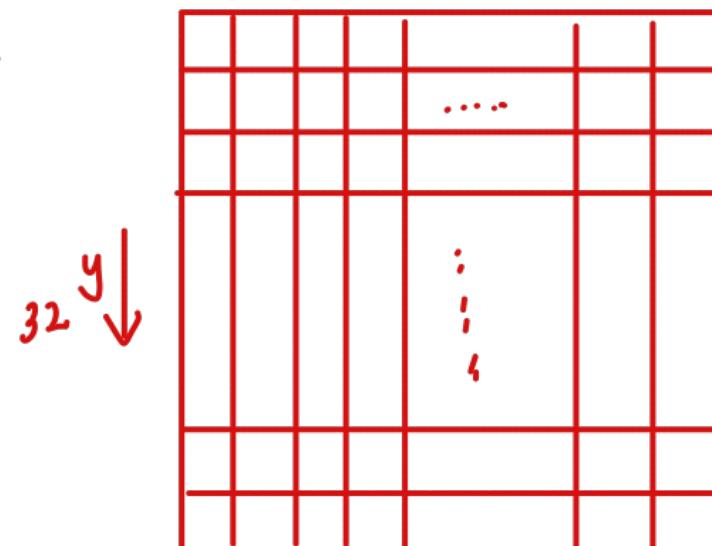
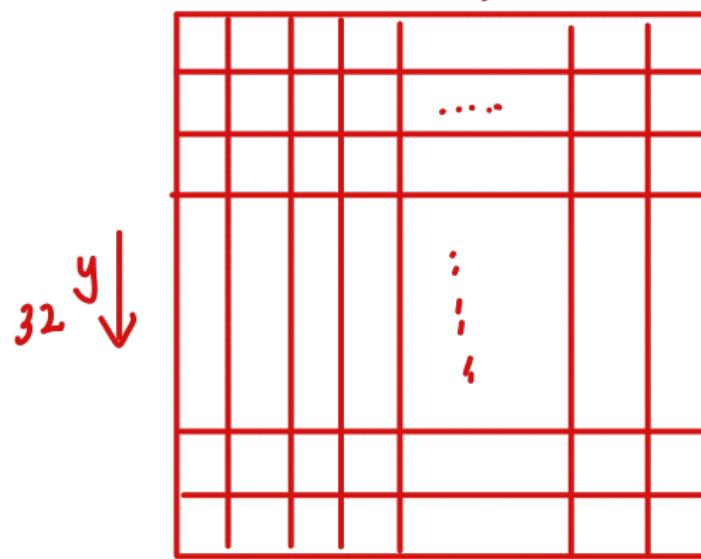
```
        int width, int height) {
```

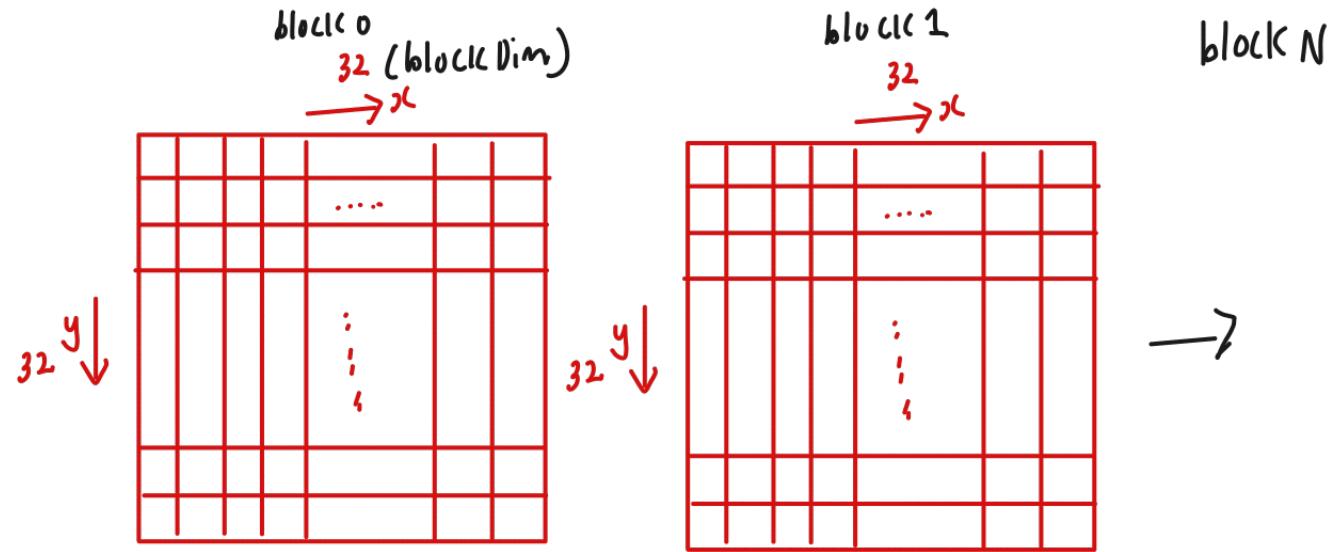
```
    int x = threadIdx.x + blockIdx.x * blockDim.x;
```

```
    int y = threadIdx.y + blockIdx.y * blockDim.y;
```

writing hot row
if ($x < width \&& y < height$) {

only process
thread inside
the img, thread
outside the img
don't run it.

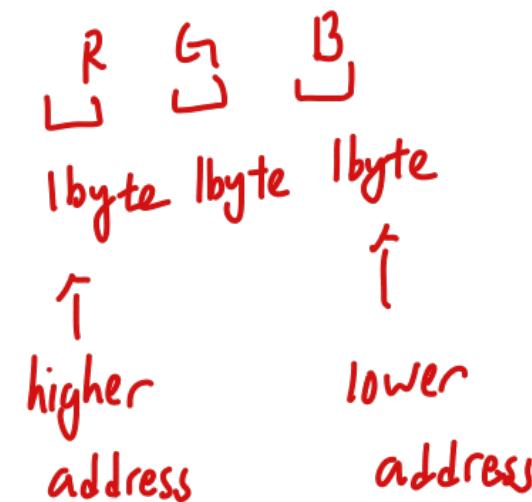




RGB TO GRayscale CONVERSION CODE

```
#define CHANNELS 3 // we have 3 channels corresponding to RGB
// The input image is encoded as unsigned characters [0, 255]
__global__ void colorConvert(unsigned char * grayImage,
                             unsigned char * rgbImage,
                             int width, int height) {
    int x = threadIdx.x + blockIdx.x * blockDim.x;
    int y = threadIdx.y + blockIdx.y * blockDim.y;

    if (x < width && y < height) {
        // get 1D coordinate for the grayscale image
        int grayOffset = y*width + x;
        // one can think of the RGB image having
        // CHANNEL times columns than the gray scale image
        int rgbOffset = grayOffset*CHANNELS;
        unsigned char r = rgbImage[rgbOffset+2]; // red value for pixel
        unsigned char g = rgbImage[rgbOffset + 1]; // green value for pixel
        unsigned char b = rgbImage[rgbOffset+2]; // blue value for pixel
    }
}
```



RGB TO GRayscale CONVERSION CODE

we abstract
think 2D



w ↪

→ x

↓ y

but ↓
Memory is linear



w w }

```
// we have 3 channels corresponding to RGB
// The input image is encoded as unsigned characters [0, 255]
__global__ void colorConvert(unsigned char * grayImage,
                             output img
                             unsigned char * rgbImage,
                             int width, int height) {
    int x = threadIdx.x + blockIdx.x * blockDim.x;
    int y = threadIdx.y + blockIdx.y * blockDim.y;

    if (x < width && y < height) {
        // get 1D coordinate for the grayscale image
        int grayOffset = y*width + x;
        // one can think of the RGB image having
        // CHANNEL times columns than the gray scale image
        int rgbOffset = grayOffset*CHANNELS;
        unsigned char r = rgbImage[rgbOffset + 3]; // red value for pixel
        unsigned char g = rgbImage[rgbOffset + 2]; // green value for pixel
        unsigned char b = rgbImage[rgbOffset + 1]; // blue value for pixel
        // perform the rescaling and store it
        // We multiply by floating point constants
        grayImage[grayOffset] = 0.21f*r + 0.71f*g + 0.07f*b;
    }
}
```

input img

(0,1) (0,2)

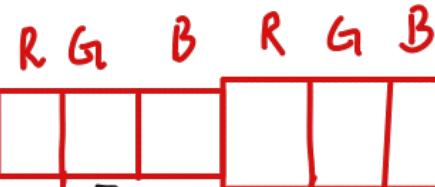
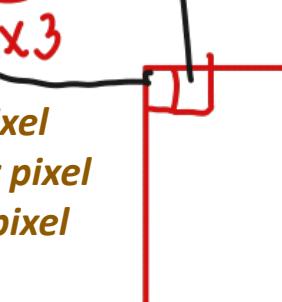
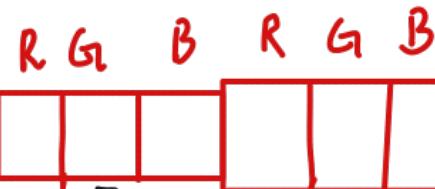


IMAGE BLURRING

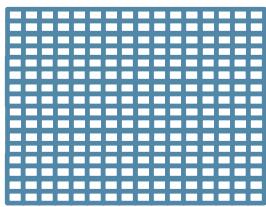


02 - Image Blurrer With Cuda

```
/**  
 * The kernel implementing the blur_image processing in CUDA  
 */  
  
__global__ void blur_image_kernel(uint8_t *blur_img, uint8_t *bgr_img,  
                                 int width, int height, int blur_size) {  
    // Compute the linear index values for x and y  
    int x = threadIdx.x + blockIdx.x * blockDim.x;  
    int y = threadIdx.y + blockIdx.y * blockDim.y;  
  
    // Blur the image. Note that each pixel is handled by one thread  
    if (x < width && y < height){  
        int pixel_idx = (y * width + x) * 3; // Each pixel has 3 components (BGR)  
  
        int sum_b = 0, sum_g = 0, sum_r = 0;  
        int num_neighbors = 0;  
  
        for (int i = -blur_size; i <= blur_size; i++){  
            for (int j = -blur_size; j <= blur_size; j++){  
                int neighbor_x = x + i;  
                int neighbor_y = y + j;  
  
                if (neighbor_x >= 0 && neighbor_x < width && neighbor_y >= 0 && neighbor_y < height){  
                    int neighbor_index = (neighbor_y * width + neighbor_x) * 3;  
  
                    sum_b += bgr_img[neighbor_index];  
                    sum_g += bgr_img[neighbor_index + 1];  
                    sum_r += bgr_img[neighbor_index + 2];  
  
                    num_neighbors++;  
                }  
            }  
        }  
  
        blur_img[pixel_idx] = sum_b / num_neighbors;  
        blur_img[pixel_idx + 1] = sum_g / num_neighbors;  
        blur_img[pixel_idx + 2] = sum_r / num_neighbors;  
    }  
}
```

```
const int thread_per_blk = 32;      // Use 32 threads per block for x and y  
void blur_image_core(uint8_t *blur_img, uint8_t *bgr_img,  
                     int width, int height, int blur_size) {  
    uint8_t* iData;           // Input image data for device  
    uint8_t* oData;           // Output image data for device  
  
    // Compute the image data size (hint: use width and height; note that  
    // the pixel size is 1 byte, and remember that each pixel has 3 colors)  
    const size_t dataSize = width * height * 3 * sizeof(uint8_t);           // Modify to the correct size  
  
    // Allocate the device memory for iData and oData the size of `dataSize`  
    cudaMalloc((void**)&iData, dataSize);  
    cudaMalloc((void**)&oData, dataSize);  
  
    // Transfer BGR image from host to device memory.  
    cudaMemcpy(iData, bgr_img, dataSize, cudaMemcpyHostToDevice);  
  
    // Calculate the blocksize and gridSize using the given `thread_per_blk`.  
    dim3 blockSize(thread_per_blk, thread_per_blk);  
    dim3 gridSize((width + thread_per_blk - 1) / thread_per_blk,  
                  (height + thread_per_blk - 1) / thread_per_blk);  
  
    // Call the CUDA blur image kernel  
    blur_image_kernel <<<gridSize, blockSize>>> (oData, iData, width, height, blur_size);  
  
    // Transfer the output (processed) image from device to host memory.  
    cudaMemcpy(blur_img, oData, dataSize, cudaMemcpyDeviceToHost);  
  
    // Free all the device memory  
    cudaFree(iData);  
    cudaFree(oData);  
}
```

BLURRING BOX



Pixels
processed by
a thread
block

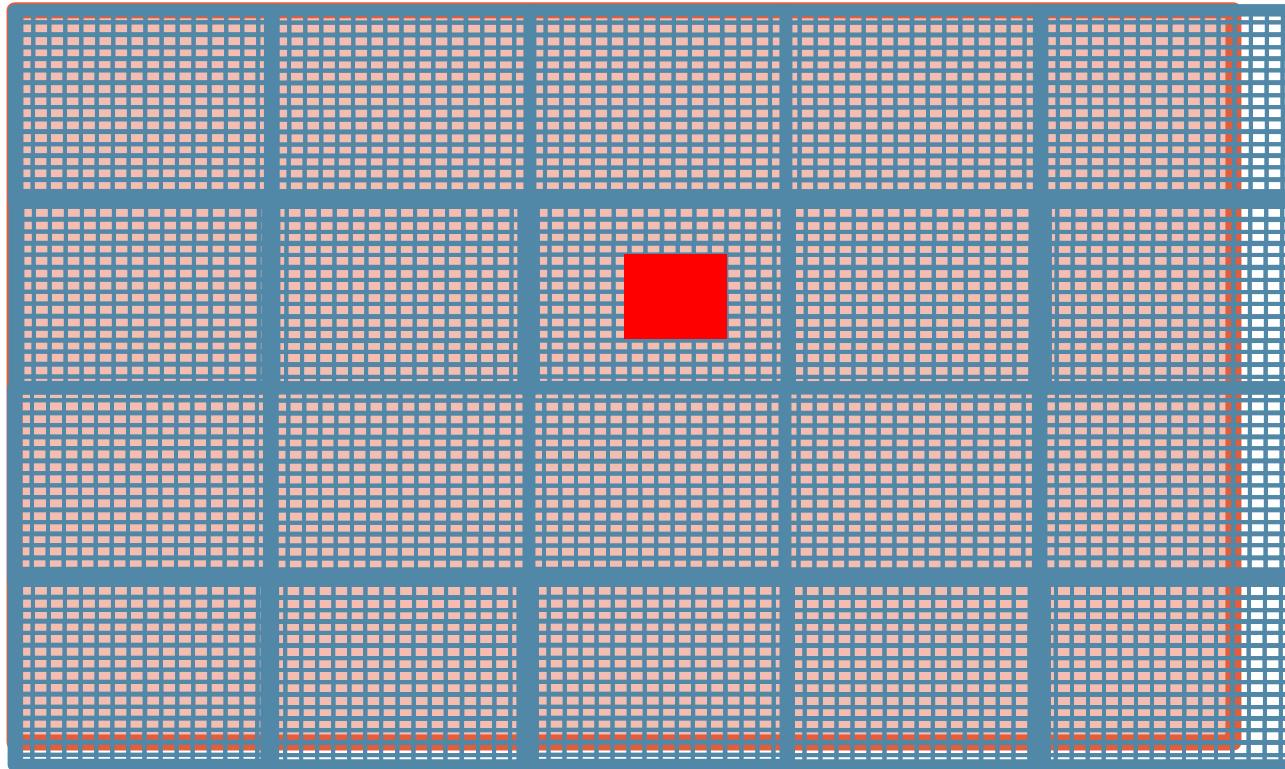


IMAGE BLUR AS A 2D KERNEL

```
__global__
void blurKernel(unsigned char * in, unsigned char * out, int w, int h)
{
    int Col = blockIdx.x * blockDim.x + threadIdx.x;
    int Row = blockIdx.y * blockDim.y + threadIdx.y;

    if (Col < w && Row < h) {
        ... // Rest of our kernel
    }
}
```

__global__

```
void blurKernel(unsigned char * in, unsigned char * out, int w, int h) {
    int Col = blockIdx.x * blockDim.x + threadIdx.x;
    int Row = blockIdx.y * blockDim.y + threadIdx.y;

    if (Col < w && Row < h) {
        int pixVal = 0;
        int pixels = 0;

        // Get the average of the surrounding 2xBLUR_SIZE x 2xBLUR_SIZE box
        for(int blurRow = -BLUR_SIZE; blurRow < BLUR_SIZE+1; ++blurRow) {
            for(int blurCol = -BLUR_SIZE; blurCol < BLUR_SIZE+1; ++blurCol) {

                int curRow = Row + blurRow;
                int curCol = Col + blurCol;
                // Verify we have a valid image pixel
                if(curRow > -1 && curRow < h && curCol > -1 && curCol < w) {
                    pixVal += in[curRow * w + curCol];
                    pixels++; // Keep track of number of pixels in the accumulated total
                }
            }
        }

        // Write our new pixel value out
        out[Row * w + Col] = (unsigned char)(pixVal / pixels);
    }
}
```

```

int curRow = curRow + blurCol;
// Verify we have a valid image pixel
if(curRow > -1 && curRow < h && curCol > -1 && curCol < w) {
    pixVal += in[curRow * w + curCol];
    pixels++; // Keep track of number of pixels in the accumulated total
}
}

// Write our new pixel value out
out[Row * w + Col] = (unsigned char)(pixVal / pixels);

```

decide how many
pixel needed to divide (in this case,
width
4 valid pixel)

add up red
4 valid pixel

add up red
25 valid pixel

do the same thing for blue
and green ↗ RGB

x	x	x
x	x	x
x	x	x

(-2, -2)

↖2	↖1	↑-2	
		↑-1	
		X	

assume BLUR_SIZE = 2

for(int blurRow = -BLUR_SIZE; blurRow < BLUR_SIZE+1; ++blurRow) {
 for(int blurCol = -BLUR_SIZE; blurCol < BLUR_SIZE+1; ++blurCol) {

blur size 1

1		
	0	
		1

blur size 2

1		
	0	
		1

↑-2

3

or

≤ BLUR_SIZE

2